

Lab 1 Performance Analysis Report

Course: CS5130

Student: Jingjing Lin

Assignment: Lab 1 — Interactive Image Mosaic Generator Using Gradio

Date: September 21, 2024

Abstract

This report presents a performance analysis of an Interactive Image Mosaic Generator that reconstructs input images using tiles from the CIFAR-10 dataset. The system employs LAB color-space matching with Euclidean (L2) distance, fully vectorized NumPy operations, and a real-time Gradio interface. Empirical results on an Apple Silicon MacBook Pro show sub-second end-to-end processing and a **71.77×** speedup versus a loop-based baseline, demonstrating suitability for interactive use.

Keywords: image mosaic, CIFAR-10, LAB color space, vectorization, NumPy, Gradio, performance analysis

1. Introduction

The Interactive Image Mosaic Generator transforms an input image into an artistic mosaic by replacing each grid cell with a semantically similar tile from a curated subset of CIFAR-10. This report evaluates computational efficiency, output quality, and scalability, and provides configuration guidance for interactive deployments.

1.1 Objectives

- Quantify runtime performance across grid resolutions.
- Compare vectorized versus loop-based implementations.
- Assess reconstruction quality using MSE, SSIM, and PSNR.
- Analyze scalability with respect to grid resolution and tile corpus size.

1.2 System Overview

- **Dataset:** CIFAR-10 (5,000 selected tiles across 10 classes)
- **Matching Space:** LAB with L2 (Euclidean) distance
- **Computation:** Fully vectorized NumPy operations (no nested Python loops)
- **Interface:** Real-time Gradio web application

2. Methodology

2.1 Processing Pipeline

1. **Image Preprocessing**
Resize to ensure divisibility by the grid; optional per-channel color quantization (16 levels); normalize to RGB.
2. **Grid Partitioning**
Vectorized reshape, e.g., `image.reshape(grid_size, cell_h, grid_size, cell_w, 3)`, to eliminate nested loops.
3. **Color Analysis**
Compute per-cell mean RGB; convert all cell means to LAB in batch for perceptual uniformity.
4. **Tile Matching**
Compute pairwise distances via broadcasting:
`np.linalg.norm(cell_colors[:, None, :] - tile_colors[None, :, :], axis=2)`; select argmin per cell.
5. **Mosaic Reconstruction**
Resize chosen tiles to cell dimensions and assemble into the final mosaic using vectorized operations.

2.2 Metrics

- **Quality:** Mean Squared Error (MSE), Structural Similarity Index (SSIM), Peak Signal-to-Noise Ratio (PSNR).
- **Computation:** End-to-end processing time (per request), initialization time (one-time), and memory footprint.

3. Experimental Setup

- **Hardware:** MacBook Pro (Apple Silicon)
- **Tile Corpus:** 5,000 CIFAR-10 images (≈ 500 per class)
- **Initialization Time:** 0.55 s (dataset load + preprocessing)
- **Test Input:** 256×256 synthetic gradient image

4. Results

4.1 Grid Resolution Study

Grid Size	Processing Time (s)	MSE	SSIM	PSNR (dB)	Total Cells
8×8	0.0135	2413.31	0.187	14.30	64
16×16	0.0493	2231.26	0.104	14.65	256
32×32	0.3854	1934.46	0.074	15.27	1024
64×64	0.7797	1505.54	0.103	16.35	4096

Findings: Processing time scales roughly with the number of cells ($O(n^2)$ in grid size). Higher resolution lowers MSE and modestly improves PSNR. A 32×32 grid offers a strong quality-latency trade-off for interactive use.

4.2 Vectorization vs. Loops

Implementation	Processing Time (s)	MSE	Speedup
Vectorized	0.0397	2231.26	71.77×
Loop-based	2.8499	2231.26	1.00×

Observation: Vectorization yields identical reconstruction quality with a 71.77× runtime reduction, enabling real-time responsiveness.

4.3 Quality Summary

- **MSE:** ~1505–2413 (decreases with resolution)
- **SSIM:** 0.074–0.187 (moderate structural similarity)
- **PSNR:** +2.05 dB improvement from 8×8 to 64×64

5. Optimization & Scalability

5.1 Computational Efficiency

- **Precomputed LAB Means:** Compute tile LAB colors once at startup.
- **View-based Ops:** Prefer reshapes/views to minimize allocations.
- **Batching/Broadcasting:** Replace Python loops with NumPy broadcasting in distance calculations.
- **Advanced Indexing:** Use argmin indices to gather tiles efficiently.

5.2 Complexity Analysis

- **Grid Resolution:** $O(n^2)$ cells.
- **Tile Corpus:** $O(m)$ with respect to number of tiles.
- **Matching:** $O(n^2 \cdot m)$ pairwise distance computation.

5.3 Practical Performance Targets

- **Interactive budget:** $\lesssim 100$ ms for highly responsive UIs (tight target).
 - **Measured:** ~ 13.5 ms (8×8) $\rightarrow \sim 779.7$ ms (64×64).
 - **Recommendation:** 32×32 grid (≈ 385 ms) balances quality and latency for typical interactive apps.
-

6. System Architecture Performance

6.1 Initialization (One-time)

- CIFAR-10 loading and preprocessing (5,000 tiles): ~ 0.55 s.
- Vectorized LAB conversion for tile means.
- **Memory footprint:** ~ 200 MB (tiles + precomputed features).

6.2 Per-Request Runtime

- **Preprocessing:** < 1 ms (resize/format).
 - **Grid Partitioning:** ~ 1 ms (reshape).
 - **Tile Matching:** 13–780 ms (grid-dependent).
 - **Reconstruction:** ~ 3 ms (tile placement/assembly).
-

7. Conclusions and Recommendations

7.1 Summary

The system achieves sub-second mosaicking for moderate grid sizes, with a **71.77 \times** acceleration from vectorization while preserving reconstruction quality (identical MSE to baseline).

LAB-space matching provides perceptually aligned selection, and the architecture scales well to large tile corpora under interactive workloads.